

Aleksander Molak

What should I buy next?

How to leverage word embeddings to build an efficient recommender system.



About me

Aleksander Molak

<https://www.linkedin.com/in/aleksandermolak/>

- Innovation Lead & Researcher at **Lingaro**
- **NLP**, sequence models, **causal modeling**
- Projects for Fortune Global 100 companies
- Complex systems, psychology, neuroscience
- Running, vegan food, languages





Overview

Recommender systems



- **Embeddings – refresher**
- **Recommender systems – refresher**
- **Recommendations with embeddings**
- **How to tune your embeddings?**

Word embeddings

Word embeddings



*An **embedding** is a relatively low-dimensional space into which you can translate high-dimensional vectors.*

Word embeddings



*An **embedding** is a relatively low-dimensional space into which you can translate high-dimensional vectors.*

REPRESENTATION



Types of embeddings

Word embeddings



- **Static**

Word embeddings

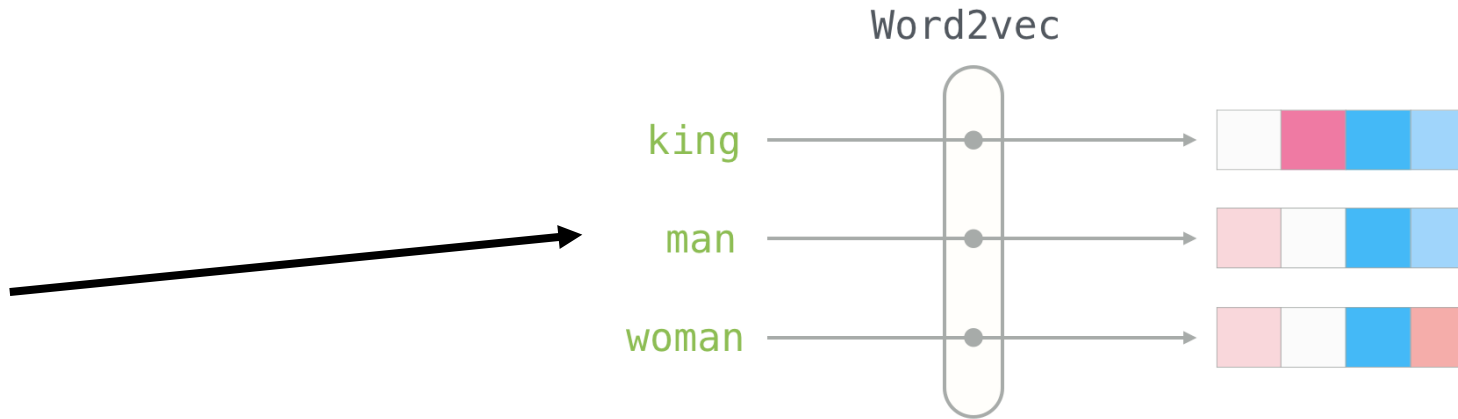


- **Static**
- **Dynamic**

Word embeddings




- **Static**
- **Dynamic**



Word embeddings



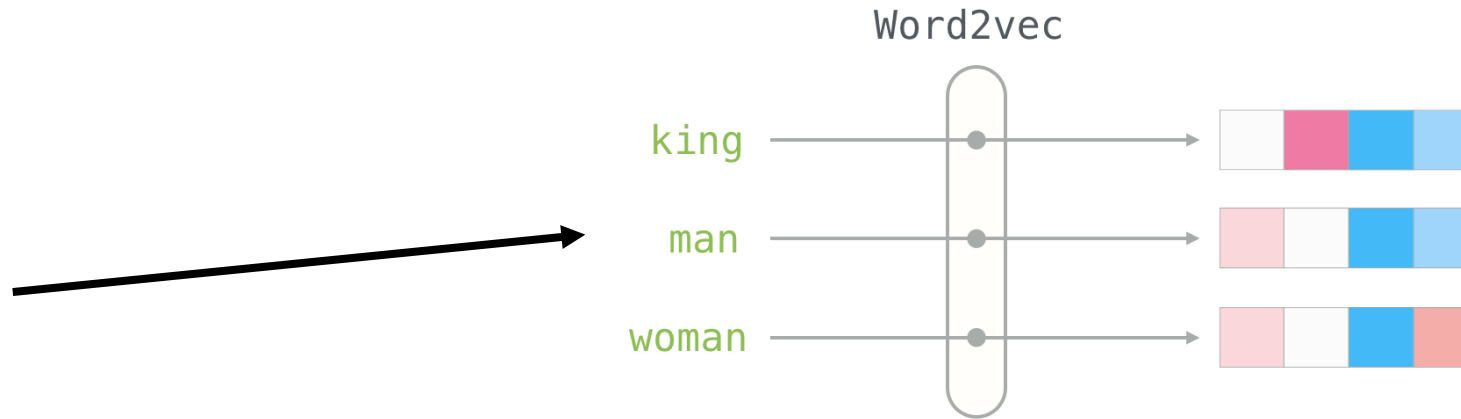
- **Static**
- **Dynamic**



Word embeddings



- Static
- Dynamic





1. I went to the bank to withdraw some cash \$
2. I spent a lovely day on the river bank! 🏞️

Word embeddings




- **Static**
- **Dynamic**



1. I went to the **bank** to withdraw some cash 💵
2. I spent a lovely day on the river **bank**! 🏞️


Word embeddings

- Static – flavors




Word embeddings

- **Static – flavors**
 - Word2Vec
 - FastText
 - Glove



Word embeddings

- **Static – flavors**
 - Word2Vec
 - FastText
 - Glove





Efficient Estimation of Word Representations in Vector Space

Tomas Mikolov

Google Inc., Mountain View, CA
tmikolov@google.com

Kai Chen

Google Inc., Mountain View, CA
kaichen@google.com

Greg Corrado

Google Inc., Mountain View, CA
gcorrado@google.com

Jeffrey Dean

Google Inc., Mountain View, CA
jeff@google.com

Word embeddings

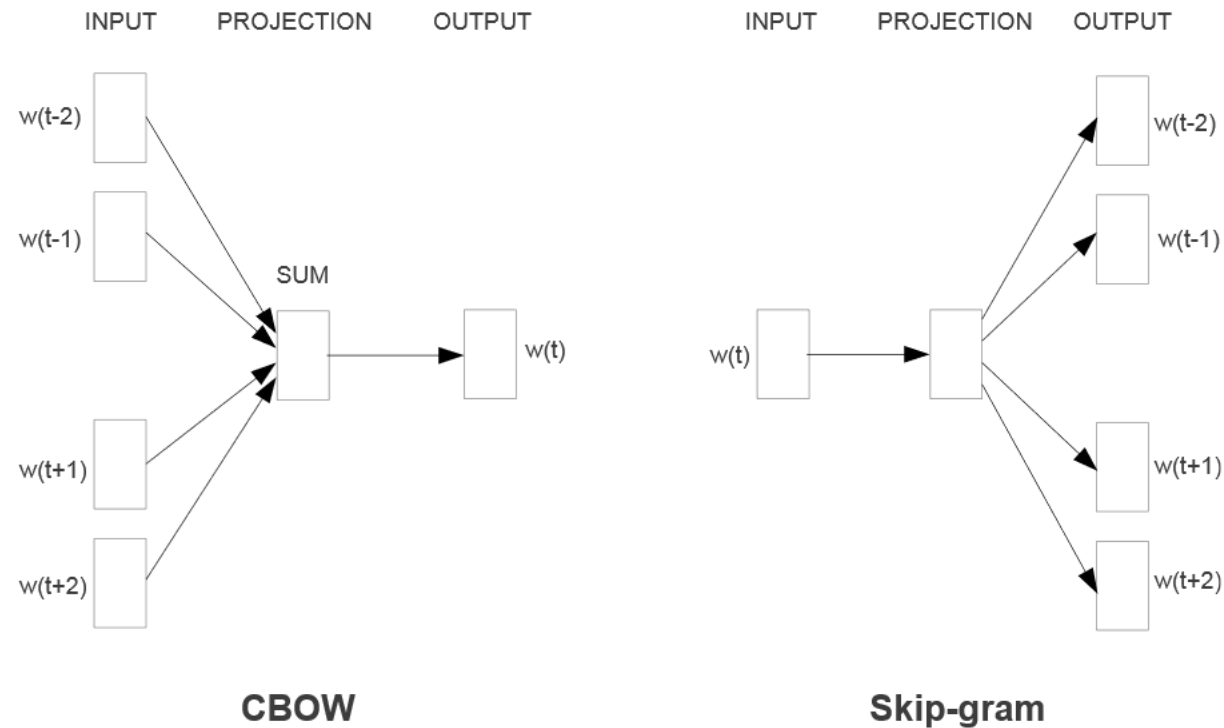


Figure 1: New model architectures. The CBOW architecture predicts the current word based on the context, and the Skip-gram predicts surrounding words given the current word.

Word embeddings

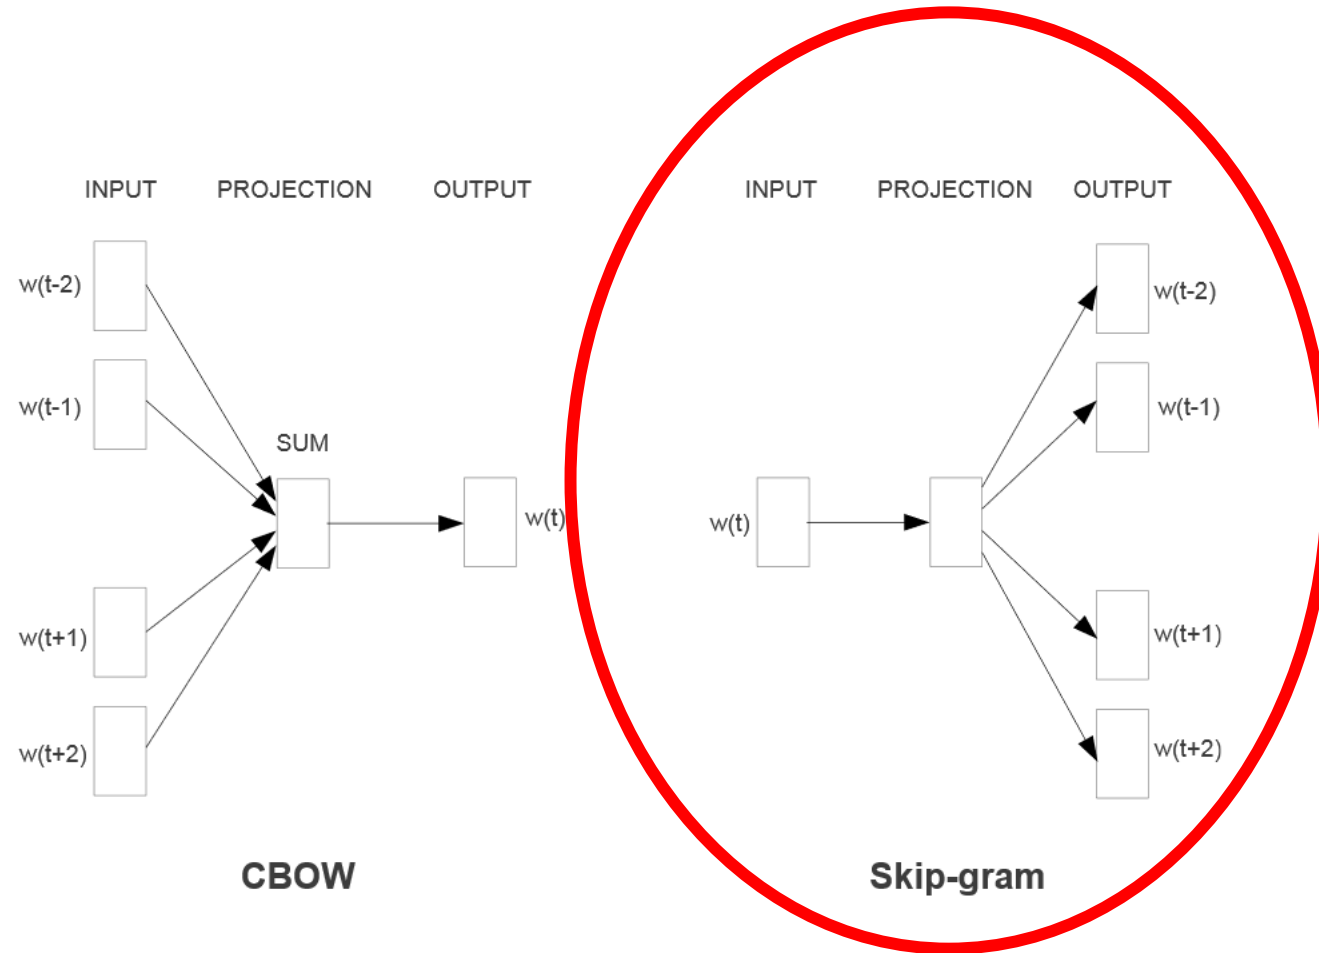



Figure 1: New model architectures. The CBOW architecture predicts the current word based on the context, and the Skip-gram predicts surrounding words given the current word.


Word embeddings



Word embeddings

***You shall know a word by the
company it keeps.***


John Rupert Firth




Word embeddings

For a large class of cases – though not for all – in which we employ the word “meaning” it can be defined thus: the meaning of a word is its use in the language.

Ludwig Wittgenstein



Word embeddings





Recommender systems

Recommender systems



Recommender systems



- Collaborative filtering



Recommender systems



- **Collaborative filtering**
- **Content-based**



Recommender systems




- Collaborative filtering
- Content-based



Recommendations with embeddings


Recommendations with embeddings

- Static embeddings – flavors




Recommendations with embeddings

- **Static embeddings – flavors**
 - Word2Vec
 - FastText
 - Glove




Recommendations with embeddings

- Static embeddings – flavors
 - Word2Vec
 - FastText
 - Glove



Recommendations with embeddings

- **Word2Vec – training regimes**
 - Word2Vec
 - MetaProd2Vec



Recommendations with embeddings




Figure 1: Prod2Vec Neural Net Architecture.

Recommendations with embeddings




Figure 1: Prod2Vec Neural Net Architecture.




Figure 2: Meta-Prod2Vec Neural Net Architecture.



Implementation

Recommendations with embeddings



```
from gensim.models import Word2Vec  
from gensim.models import KeyedVectors
```

Recommendations with embeddings



```
# Train the model  
model = Word2Vec(train_data, **model_params, callbacks=callbacks)
```



Word2Vec vs Meta-Prod2Vec





Hyperparameters

Recommendations with embeddings



Sub-sampling


Sub-sampling reduces probability of using the most frequent (or the least frequent) words in the training.

Note that the formula comes from Google implementation of W2V and it's slightly different from the one in the paper.

$$P(w_j) = \left(\sqrt{\frac{z(w_j)}{k}} + 1 \right) \frac{k}{z(w_j)}$$

** Default value for $k = .001$

Recommendations with embeddings



Recommendations with embeddings




Negative sampling

Instead of all negative examples (all words not present in the context), we will just update weights for a couple of randomly selected words.

$$P(w_i) = \frac{f(w_i)^k}{\sum_{j=0}^N (f(w_j)^k)}$$

Recommendations with embeddings



Recommendations with embeddings



Word2vec applied to Recommendation: Hyperparameters Matter

Hugo Caselles-Dupré^{12*}

¹ Flowers Laboratory (ENSTA
ParisTech & INRIA)

² Softbank Robotics Europe
Paris, France
caselles@ensta.fr

Florian Lesaint

Deezer SA
Paris, France
flesaint@deezer.com

Jimena Royo-Letelier

Deezer SA
Paris, France
jroyo@deezer.com

Recommendations with embeddings



```
# Initialize callbacks
callbacks = [EpochLogger()]

# Set model params
EMB_DIM = 300
N_CORES = 4
WINDOW = 9    # For traditional w2v you might want to try the 90th percentile of len(rows)
N_EPOCHS = 200
NEG_EXP = -.5
MIN_COUNT = 5
SAMPLE = 0

model_params = dict(size = EMB_DIM,
                    sg = 1,
                    negative = 5,
                    iter = N_EPOCHS,
                    ns_exponent = NEG_EXP,
                    sample = SAMPLE,
                    workers = N_CORES,
                    window = WINDOW,
                    min_count = MIN_COUNT)

# Train the model
model = Word2Vec(train_data, **model_params, callbacks=callbacks)
```

Pros and cons

Pros



- **Easy to maintain**
- **Supports continuous training**
- **Fast to train**
- **Ultra fast retrieval (e.g. with HNSW*)**
- **Can be easily served in a serverless fashion** (depending on the no. of tokens)

* Paper: <https://arxiv.org/ftp/arxiv/papers/1603/1603.09320.pdf>

C++ / Python implementation: <https://github.com/nmslib/hnswlib>

Cons 🙄

- **Not suitable for cold-start scenarios**
- **Might be noisy**
- **Not easily interpretable**



Summary

Summary



- **Embedding-based systems are fast and easily maintainable**
- **Word2Vec and Meta-Prod2Vec cover a broad spectrum of cases**
- **HNSW gives you an amazing recall / speed ratio**
- **It's worth to pay attention to hyperparams if you want good results**

Thank you!

Thank you!

Aleksander Molak

LinkedIn

<https://www.linkedin.com/in/aleksandermolak/>

Twitter

@AleksanderMolak

